

Peak2Cloud: Scientific Computing on the Cloud

JOSEPH ANTHONY C. HERMOCILLA, University of the Philippines Los Banos

Peak2Cloud (P2C) is an Openstack-based private cloud for scientific and high performance computing. First, we present how P2C was configured and tested. Then we describe vcluster, a tool for rapidly deploying message-passing clusters on P2C. Lastly, we analyze some benchmark results on the performance of P2C deployed virtual clusters.

Categories and Subject Descriptors: C.2.4 [Computer-Communication Networks]: Distributed Systems

General Terms: Network operating systems

Additional Key Words and Phrases: cloud computing, high-performance computing

ACM Reference Format:

Joseph Anthony C. Hermocilla, 2010. Peak2Cloud: Scientific Computing on the Cloud. *ACM Trans. Comput. Syst.* 2014, 1, Article 2 (June 2014), 2 pages.

DOI: <http://dx.doi.org/10.1145/0000000.0000000>

1. INTRODUCTION

Cloud computing has become a buzzword in today's modern computing, though there is no agreed upon meaning of the term. In 2011, NIST [Mell and Grance 2011] published a definition that is widely quoted and used. The popularity of cloud computing mainly comes from its ability to provision additional resources on demand with minimum intervention from the provider. It leverages advances in virtualization and web services technologies. For example, a website with a sudden increase in workload can start another server machine (virtual) almost instantaneously to accommodate the additional load.

Cloud computing offers service models which include Software-as-a-Service(SaaS), Platform-as-a-Service(PaaS), and Infrastructure-as-a-Service(IaaS). IaaS allows the consumer to provision computing resources(hardware, network, storage) to run arbitrary software including operating systems [Mell and Grance 2011].

2. RELATED WORK

Studies have been published to evaluate the applicability of the cloud for scientific computing.[Ekanayake and Fox 2010] [Evangelinos and Hill 2008] [Expsito et al. 2013][Ludescher et al. 2013] [Mauch et al. 2013][Jackson et al. 2010][Zhai et al. 2011] [Walker].Most of these utilized the public cloud, specifically Amazon EC2 as their testbed.

This work is supported by the Philippine Department of Science and Technology (DOST) Accelerated Science and Technology Human Resource Development Program (ASTHRDP).

Author's address: J. A. C. Hermocilla, Institute of Computer Science, University of the Philippines Los Banos

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or permissions@acm.org.

© 2014 ACM 0734-2071/2014/06-ART2 \$15.00

DOI: <http://dx.doi.org/10.1145/0000000.0000000>

3. METHODOLOGY

3.1. Openstack

Openstack an open source software framework for deploying clouds. It is based on Nebula used by NASA. It provides a public interface that is compatible with Amazon EC2. There are several components of Openstack which are developed separately. Keystone is used to manage authorization and authentication. Glance manages the virtual machine images. The Nova component is used for managing compute nodes.

3.2. Hardware

3.3. Network Topology

3.4. vcluster

3.5. Benchmarks

4. RESULTS AND DISCUSSION

5. CONCLUSIONS

ACKNOWLEDGMENTS

The author would like to thank the Lord.

REFERENCES

- Jaliya Ekanayake and Geoffrey Fox. 2010. High performance parallel computing with clouds and cloud technologies. In *Cloud Computing*. Springer, 2038. http://link.springer.com/chapter/10.1007/978-3-642-12636-9_2
- Constantinos Evangelinos and C. Hill. 2008. Cloud computing for parallel scientific hpc applications: Feasibility of running coupled atmosphere-ocean climate models on amazons ec2. *ratio* 2, 2.40 (2008), 234. <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.296.3779&rep=rep1&type=pdf>
- Roberto R. Expsito, Guillermo L. Taboada, Sabela Ramos, Juan Tourio, and Ramn Doallo. 2013. Performance analysis of HPC applications in the cloud. *Future Generation Computer Systems* 29, 1 (Jan. 2013), 218–229. DOI: <http://dx.doi.org/10.1016/j.future.2012.06.009>
- Keith R. Jackson, Lavanya Ramakrishnan, Krishna Muriki, Shane Canon, Shreyas Cholia, John Shalf, Harvey J. Wasserman, and Nicholas J. Wright. 2010. Performance Analysis of High Performance Computing Applications on the Amazon Web Services Cloud. *IEEE*, 159–168. DOI: <http://dx.doi.org/10.1109/CloudCom.2010.69>
- Thomas Ludescher, Thomas Feilhauer, and Peter Brezany. 2013. Cloud-Based Code Execution Framework for scientific problem solving environments. *Journal of Cloud Computing: Advances, Systems and Applications* 2, 1 (2013), 11. <http://www.journalofcloudcomputing.com/content/2/1/11>
- Viktor Mauch, Marcel Kunze, and Marius Hillenbrand. 2013. High performance cloud computing. *Including Special sections: High Performance Computing in the Cloud & Resource Discovery Mechanisms for P2P Systems* 29, 6 (Aug. 2013), 1408–1416. DOI: <http://dx.doi.org/10.1016/j.future.2012.03.011>
- Peter Mell and Timothy Grance. 2011. The NIST definition of cloud computing (draft). *NIST special publication* 800, 145 (2011), 7. <http://pre-developer.att.com/home/learn/enablingtechnologies/The.NIST.Definition.of.Cloud.Computing.pdf>
- Edward Walker. Benchmarking Amazon EC2 for high-performance scientific computing. (????). <https://www.usenix.org/legacy/publications/login/2008-10/openpdfs/walker.pdf>
- Yan Zhai, Mingliang Liu, Jidong Zhai, Xiaosong Ma, and Wenguang Chen. 2011. Cloud versus in-house cluster: evaluating Amazon cluster compute instances for running MPI applications. In *State of the Practice Reports*. ACM, 11. <http://dl.acm.org/citation.cfm?id=2063363>

Received May 2014; revised June 2014; accepted June 2014